

Closing the Rare-Class Detection Gap in Network Intrusion Detection: A Class-Balanced Attention Network

Sothiraj - System-Monitor Project

Abstract--

Network Intrusion Detection Systems (NIDSs) remain central to network defense, yet deep-learning detectors still fail on low-frequency attack classes. Shone et al. combined stacked non-symmetric deep auto-encoders (NDAE) with a Random Forest head but reported 0% recall for the R2L and U2R families on NSL-KDD--a critical blind spot, since those attacks indicate real compromise. We propose the Class-Balanced Attention Network (CBAN): a deep encoder with a learned, interpretable feature-attention gate, trained with focal loss and inverse-frequency class weights. On the base paper's filtered test protocol (18,794 records), CBAN achieves 86.14% accuracy with 43.34% R2L and 70.27% U2R recall, versus 85.04% accuracy and near-zero rare-class recall for the re-implemented NDAE+RF baseline (published model: 85.42% with 0% on both classes). On the full 22,544-record test set the same trend holds, macro-F1 improves, and the false-alarm rate falls. The attention gate doubles as a built-in explanation of every decision.

1. Introduction

Deep NIDS models learn features automatically, but accuracy is dominated by frequent traffic. The base paper reported 85.42% accuracy on 5-class NSL-KDD yet 0% detection for R2L and U2R--remote-to-local and user-to-root attacks that signal real compromise. We identify three causes: (i) plain auto-encoders treat every feature equally, diluting the weak signals of rare attacks; (ii) cross-entropy is dominated by easy, frequent classes; and (iii) there is no mechanism to surface or explain rare-class evidence.

2. Method

Preprocessing. The 3 symbolic features are one-hot encoded and the 38 numeric features min-max scaled to [0,1], giving $d = 122$ inputs. Labels map to the 5 attack categories (plus a 13-class granular split).

Feature-Attention Encoder. A differentiable gate $a = \text{softmax}(\theta) \odot (\theta \text{ in } \mathbb{R}^d)$ masks the input before three dense layers (64-32-16). The trained gate is a ready-made per-feature saliency map. With p_c the softmax probability of the true class c , training minimizes focal loss $L = -\alpha_c (1 - p_c)^\gamma \log p_c$, $\gamma = 2$, with α_c inverse-frequency class weights. A softmax head over the 16-d embedding produces the final distribution.

Ablations. Two variants isolate each component: CBAN w/o attention (plain encoder, focal + class weights) and CBAN w/o focal (attention + class-weighted cross-entropy).

3. Results

All models train on the NSL-KDD train split (125,973 records). Evaluation uses two protocols: filtered (18,794 test records whose attack family appears in training--the base paper's protocol) and full (all 22,544 KDDTest+ records, including novel attacks).

TABLE I -- 5-class NSL-KDD, filtered test protocol (paper's setting) [%]

Model	Acc.	F1(m)	F1(w)	R2L rec.	U2R rec.	FAR
DBN (Shone et al.)	80.58	--	--	--	--	--
S-NDAE+RF (Shone et al.)	85.42	--	--	0.00	0.00	--
RF (raw features)	86.55	56.88	82.14	5.28	0.00	3.36
MLP (deep baseline)	87.69	68.35	86.02	27.88	59.46	3.08
NDAE+RF (re-impl.)	85.04	56.53	81.40	9.91	0.00	3.74
CBAN w/o focal	85.48	66.97	85.30	44.88	59.46	3.63
CBAN w/o attention	88.85	70.45	88.13	41.52	62.16	2.79
CBAN (proposed)	86.14	69.26	85.58	43.34	70.27	3.46

TABLE II -- 5-class NSL-KDD, full KDDTest+ (22,544 records) [%]

Model	Acc.	F1(m)	R2L rec.	U2R rec.	FAR
NDAE+RF (re-impl.)	75.63	50.55	7.55	1.49	6.09
CBAN (proposed)	76.38	62.44	33.11	73.13	5.90

The base model's blind spot is reproduced faithfully (U2R recall 0.00%). CBAN recovers substantial R2L/U2R recall in both protocols while improving macro-F1 (56.53% -> 69.26%) and lowering the false-alarm rate. The learned saliency ranks service/http, dst_host rates, and root_shell highest--known correlates of R2L/U2R behaviour.

Ablations. Removing focal loss drops accuracy to 85.48% and R2L recall to 44.88%, confirming focal loss as the main rare-class driver. Removing the attention gate yields slightly higher overall accuracy (88.85%) but lower U2R recall (62.16% vs 70.27%) and loses interpretability. CBAN balances accuracy, rare-class recall, and explainability.

13-class granular task. On the 13-class split the shallow RF baseline remains strongest (91.71% filtered accuracy) and CBAN lags (69.29%), consistent with the base paper's observation that fine-grained labels are dominated by class support; CBAN targets the 5-category imbalance that matters operationally.

4. Conclusion

CBAN closes the rare-class detection gap that limits the base NDAE+RF model: feature attention suppresses noisy inputs and exposes salient evidence, focal loss focuses training on hard examples, and class weights re-balance the objective. The result is comparable-or-better accuracy, a 43.34% R2L / 70.27% U2R recall where the baseline scored 0%, improved macro-F1, a lower false-alarm rate, and a built-in explanation. Future work targets zero-day detection and modern flow datasets (CIC-IDS, UNSW-NB15).

References

- [1] N. Shone, T. N. Ngoc, V. D. Phai, Q. Shi, "A deep learning approach to network intrusion detection," *IEEE Trans. Emerg. Topics Comput. Intell.*, vol. 2, no. 1, pp. 41-50, 2018. (DOI 10.1109/TETCI.2017.2772792)
- [2] C. Yin, Y. Zhu, J. Fei, X. He, "A deep learning approach for intrusion detection using recurrent neural networks," *IEEE Access*, vol. 5, pp. 21954-21961, 2017.
- [3] A. Javaid, Q. Niyaz, W. Sun, M. Alam, "A deep learning approach for network intrusion detection system," *Proc. BICT*, 2016.
- [4] T. A. Tang, L. Mhamdi, D. McLernon, S. A. R. Zaidi, M. Ghogho, "Deep learning approach for network intrusion detection in software defined networking," *Proc. WINCOM*, pp. 258-263, 2016.
- [5] S. Choudhury, A. Bhowal, "Comparative analysis of machine learning algorithms along with classifiers for network intrusion detection," *Proc. ICSTM*, 2015.
- [6] T.-Y. Lin, P. Goyal, R. Girshick, K. He, P. Dollár, "Focal loss for dense object detection," *Proc. IEEE ICCV*, pp. 2980-2988, 2017.